

```
Bk
from flask import Flask, request, jsonify

from flask_cors import CORS

import openai

import numpy as np


app = Flask(__name__)

CORS(app)


# -----
# Azure OpenAI Configuration
# -----

openai.api_type = "azure"

openai.api_base = "https://hemal-mg3addke-swedencentral.openai.azure.com/"

openai.api_version = "2024-12-01-preview"

openai.api_key =
"70DsK0AZNsgtS6zeRdyeO3lgKX9Cci2mhXpuBPv0ezrPaRpYkOmVJQQJ99BIACfhMk5XJ3w
3AAAAACOGCBsh"


deployment_name = "gpt-35-turbo-demo"    # Chat model deployment name
embedding_model = "embedding-demo"      # Embedding model deployment name


# -----
# Example documents
# -----

documents = [

    {"id": 1, "text": "Angular is a frontend framework by Google."},
```

```
{ "id": 2, "text": "Flask is a lightweight Python web framework." },
{ "id": 3, "text": "Azure OpenAI service allows GPT models in the cloud." }
]
```

```
# -----
```

```
# Precompute embeddings for documents
```

```
# -----
```

```
document_embeddings = []
```

```
for doc in documents:
```

```
    emb = openai.Embedding.create(
```

```
        input=doc["text"],
```

```
        engine=embedding_model # Use 'engine' for Azure
```

```
    )["data"][0]["embedding"]
```

```
    document_embeddings.append({ "id": doc["id"], "text": doc["text"], "embedding": emb })
```

```
# -----
```

```
# Cosine similarity function
```

```
# -----
```

```
def cosine_similarity(a, b):
```

```
    a = np.array(a)
```

```
    b = np.array(b)
```

```
    return np.dot(a, b) / (np.linalg.norm(a) * np.linalg.norm(b))
```

```
# -----
```

```
# API route for frontend
```

```
# -----
```

```

@app.route("/ask", methods=["POST"])
def ask():
    data = request.json
    question = data.get("question", "")
    if not question:
        return jsonify({"answer": "No question provided."})

    # Step 1: Get embedding of the question
    q_emb = openai.Embedding.create(
        input=question,
        engine=embedding_model # Use 'engine' for Azure
    )['data'][0]['embedding']

    # Step 2: Find most similar document
    similarities = [cosine_similarity(q_emb, doc["embedding"]) for doc in
document_embeddings]

    best_index = int(np.argmax(similarities))
    top_doc = document_embeddings[best_index]["text"]

    # Step 3: Ask GPT with the top document as context
    try:
        response = openai.ChatCompletion.create(
            engine=deployment_name, # Use 'engine' for Azure
            messages=[
                {"role": "system", "content": "You are a helpful assistant."},
                {"role": "user", "content": f"Answer the question based on this document:
{top_doc}\nQuestion: {question}"}

```

```

    ],
    temperature=0.7
)
answer = response['choices'][0]['message']['content']
return jsonify({"answer": answer})

except Exception as e:
    return jsonify({"answer": f"Error: {str(e)}"})

# -----
# Run server
# -----
if __name__ == "__main__":
    app.run(debug=True)

home.ts

import { Component } from '@angular/core';
import { CommonModule } from '@angular/common';
import { FormsModule } from '@angular/forms';
import { HttpClient, HttpClientModule } from '@angular/common/http';

@Component({
  selector: 'app-home',
  standalone: true,
  imports: [CommonModule, FormsModule, HttpClientModule],
  templateUrl: './home.html',

```

```

    styleUrls: ['./home.css']
  })
}

export class Home {

  prompt: string = "";

  isSubmitting = false;

  submittedPrompts: string[] = [];

  backendResponse: string = ""; // 👉 store backend response

  constructor(private http: HttpClient) {}

  submitPrompt(): void {

    const text = this.prompt.trim();

    if (!text || this.isSubmitting) return;

    this.isSubmitting = true;

    this.backendResponse = "";

    // ✅ Call backend API instead of just simulating
    this.http.post<{ answer: string }>('http://127.0.0.1:5000/ask', { question: text })
      .subscribe({
        next: (res) => {
          console.log('Backend response:', res);

          this.submittedPrompts.push(text);

          this.backendResponse = res.answer; // store response

          this.prompt = "";

          this.isSubmitting = false;
        }
      });
  }
}

```

```

    },
    error: (err) => {
      console.error('Error:', err);

      this.backendResponse = ' ⚠️ Could not connect to backend';
      this.isSubmitting = false;
    }
  });
}

```

```

clearPrompt(): void {
  if (this.isSubmitting) return;
  this.prompt = "";
}

```

```

get isSubmitDisabled(): boolean {
  return !((this.prompt ?? "").trim()) || this.isSubmitting;
}
}

```

home.html

```

<div class="home-center dark-theme" style="--header-h: 0px;">
  <section class="prompt-card" role="region" aria-labelledby="promptTitle">
    <h2 class="title" id="promptTitle">Ask anything</h2>

    <label class="sr-only" for="promptInput">Your prompt</label>
    <textarea

```

```
id="promptInput"
class="prompt-input"
[(ngModel)]="prompt"
placeholder="Type your prompt here..."
(keydown.control.enter)="submitPrompt()"
(keydown.meta.enter)="submitPrompt()"
></textarea>

<div class="actions">

  <button class="clear-btn" type="button" (click)="clearPrompt()"
[disabled]="!prompt.length || isSubmitting">

    Clear

  </button>

  <button class="primary-btn" type="button" (click)="submitPrompt()"
[disabled]="isSubmitDisabled">

    {{ isSubmitting ? 'Submitting...' : 'Submit' }}

  </button>

</div>

<!-- Submitted prompts displayed below -->

<div class="submitted-prompts" *ngIf="submittedPrompts.length">

  <h3 class="submitted-title">Submitted Prompts:</h3>

  <div class="prompt-card-small" *ngFor="let p of submittedPrompts">

    {{ p }}

  </div>

</div>
```

```
<!-- 📢 Show backend response -->

<div *ngIf="backendResponse" class="response-box">

  <h3>Backend Response:</h3>

  <p>{{ backendResponse }}</p>

</div>

</section>

</div>
```